



Advanced Databases Course project

Team 9

Ayman Mohamed Reda

Ahmed Ayman

Ammar Mohamed

Omar Ahmed Rezk

Proposes:

1) Schema Optimizations:

we will merge 2 tables together instead of having them discrete.

we will modify column data Types, for example: from Float to Int, from BigInt to Int, from nvarchar to varchar or to char, from char to tinyInt or Bool (in gender column).

As we are trying to minimize Table sizes as much as possible.

2) Memory and Cache optimization:

We will increase Buffer Size from the DBMS.

We tried to increase Block size but we couldn't do it, DBMS didn't let us to.

Tried to store ready procedures, DBMS is far better in this task, it optimizes the query very heavily if you run the query multiple times in a row.

3) Index Tuning:

We added 2 indices, one for Salary on table Employee only and the other (Salary, SSN) on table Employee. Besides the default SSN index on table Employee.

4) Query Optimization:

We tried to get an optimized version for all our queries, we succeeded in 2 of them and saw great improvements.

Validation and Testing:

All tests were done on the same computer, same environment with the charger plugged in and zero power savings.

PC specs:

Device name	LAPTOP-66BOPJMV
Processor	Intel(R) Core(TM) i5-10300H CPU @ 2.50GHz 2.50 GHz
Installed RAM	16.0 GB (15.8 GB usable)
System type	64-bit operating system, x64-based processor

OS Details:

Edition Windows 10 Home

Version 21H2

Installed on 4/7/2021

OS build 19044.1645

Experience Windows Feature Experience Pack 120.2212.4170.0

GPU: NVIDIA GEFORCE GTX 1650

DISKS: ST1000LM048-2E7172 (1 TB)

SSD: WDC PC SN530 SDBPNPZ-256G-1014 (256GB)

Experimentation steps:

We implemented and filled 9 variants of databases, with their description below:

Type	Size	Optimization
SQL	10K	False
SQL	100K	False
SQL	1M	False
SQL	10M	False
SQL	100K	Index Tuning
SQL	100K	Increase Buffer Size + Index Tuning
SQL	100K	Optimized Schema only
SQL	100K	Full optimization (schema + index + Buffer size)
noSQL	100K	False

- For the Raw Databases, we test with normal queries.
- Then we test with optimized queries
- Then we test index-tuned DB with normal queries vs no index-tuning
- Then we test increased Buffer Sized DB vs normal conditions
- Then we test Full optimization DB vs zero optimization
- Then we test Full optimization DB vs noSQL DB

Results and Queries Follows:

Proposes:

1) Schema Optimizations:

we will merge 2 tables together instead of having them discrete.

we will modify column data Types, for example: from Float to Int, from BigInt to Int, from nvarchar to varchar or to char, from char to tinyInt or Bool (in gender column).

As we are trying to minimize Table sizes as much as possible.

2) Memory and Cache optimization:

We will increase Buffer Size from the DBMS.

We tried to increase Block size but we couldn't do it, DBMS didn't let us to.

Tried to store ready procedures, DBMS is far better in this task, it optimizes the query very heavily if you run the query multiple times in a row.

3) Index Tuning:

We added 2 indices, one for Salary on table Employee only and the other (Salary, SSN) on table Employee. Besides the default SSN index on table Employee.

4) Query Optimization:

We tried to get an optimized version for all our queries, we succeeded in 2 of them and saw great improvements.

Validation and Testing:

All tests were done on the same computer, same environment with the charger plugged in and zero power savings.

PC specs:

Device name LAPTOP-66BOPJMV
Processor Intel(R) Core(TM) i5-10300H CPU @ 2.50GHz 2.50 GHz
Installed RAM 16.0 GB (15.8 GB usable)
System type 64-bit operating system, x64-based processor

OS Details:

Edition Windows 10 Home
Version 21H2
Installed on 4/7/2021
OS build 19044.1645
Experience Windows Feature Experience Pack 120.2212.4170.0

GPU: NVIDIA GEFORCE GTX 1650

DISKS: ST1000LM048-2E7172 (1 TB)
SSD: WDC PC SN530 SDBPNPZ-256G-1014 (256GB)

Experimentation steps:

We implemented and filled 9 variants of databases, with their description below:

Type	Size	Optimization
SQL	10K	False
SQL	100K	False
SQL	1M	False
SQL	10M	False
SQL	100K	Index Tuning
SQL	100K	Increase Buffer Size + Index Tuning
SQL	100K	Optimized Schema only
SQL	100K	Full optimization (schema + index + Buffer size)
noSQL	100K	False

- For the Raw Databases, we test with normal queries.
- Then we test with optimized queries
- Then we test index-tuned DB with normal queries vs no index-tuning
- Then we test increased Buffer Sized DB vs normal conditions

- Then we test Full optimization DB vs zero optimization
- Then we test Full optimization DB vs noSQL DB

Results and Queries Follows:

```
another_complex3 = 'select count(SSN) from Employee a, Department b where
a.SSN < b.Dnumber and a.Salary < b.Dnumber and b.Mgr_SSN < 80000'
```

```
another_complex_ayman = 'select DISTINCT Salary, Fname, Dname from
Employee e, Department d, Contracts c where e.SSN = c.Emp_SSN and e.DNO =
d.Dnumber'
```

```
another_complex4 = 'select SSN , Salary from Employee a, Department b
where a.SSN < b.Dnumber and a.Salary < b.Dnumber and b.Mgr_SSN < 80000
group by Salary, SSN;'
```

Times in minutes

	Query1	Query2	Query3	Query4
10k	0.0621	0.112985	0.03771735	2.354
100k	0.7951	1.382995	4.403953	224.1677
1M	1.078	0.930993	81.8903	
10M	1.181997	7.10634		

AFTER QUERY OPTIMIZATIONS:



```
another_complex3_OPT = 'select count(SSN) from Employee a inner join
Department b on a.SSN < b.Dnumber inner join Department d on a.Salary <
d.Dnumber where d.Mgr_SSN < 80000'
```

```
another_complex4_OPT = 'select SSN, Salary from (Select SSN, Salary from
Employee)as a inner join (select Dnumber, Mgr_SSN from Department where
Mgr_SSN < 80000)as b on a.SSN < b.Dnumber group by Salary, SSN'
```

Times in minutes

	Q3 opt	Q4 opt		
10k	0.064999	0.005024		
100k	3.4	0.00506		
1M	78.58	0.00499		
10M				

Why on Q4 we can see that the query takes a constant time?

Because we selected Mgr_SSN < 80000 as early as possible, shrinking our search space to multiples of 80000 only, not multiples of millions of records.

After putting index on column Salary we will clearly see Optimizations when executing Q4

Time in seconds

	Query 4 without optimizations After Salary Index
10k	2.59904
100k	1.94004
1M	78.32
10M	102.4084717

Now we will increase the buffer size

From 65536 B to 655360 B

	Query 4 without optimizations
10k	2.26903
100k	1.86799
1M	120.9788
10M	108.405499

We will try to push it up to 65536000 B

	Query 4 without optimizations
10k	2.343
100k	1.9679
1M	107.55343

Now reverting all our optimizations again (deleting the additional indices, decreasing buffer size, no optimizations) and getting worse results again, to prove that our optimizations were effective.

Results in seconds

	Query 4
10k	3.313
100k	315.242124
1M	2490.514

Try to optimize the schema by merging Works_On with Employee, by adding columns “PNO” & “Hours” to Employee instead of having an entire different table.

And came up with a new query:

```
Same_query = 'select sum(Hours) from Employee as e, Project as p,  
Department as d where e.PNO = p.Pnumber and p.DNO = d.Dnumber and  
d.Dnumber > 50 and d.Mgr_SSN = e.SSN and e.Hours > 10 group by e.Salary,  
d.Dnumber'
```

Results in seconds:

	Before merging	After Merging
100k DB	6.09847	0.124001

Space improvement:

Employee size + Works_On size before merging = 40.352 kB

Employee size after merging = 39.291

Space improvement = 102.62%

OPTIMIZATIONS DONE:

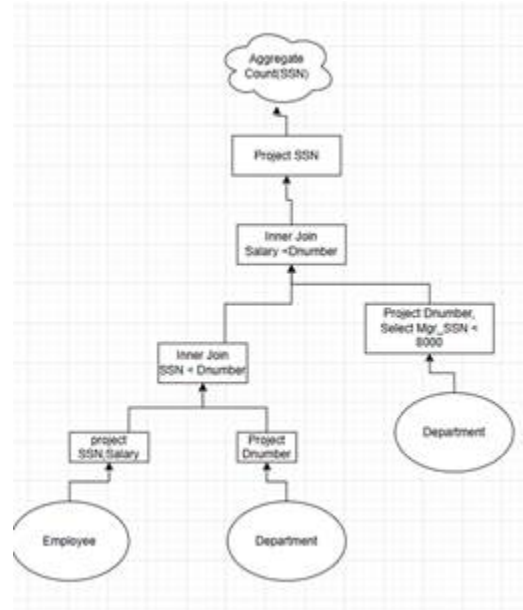
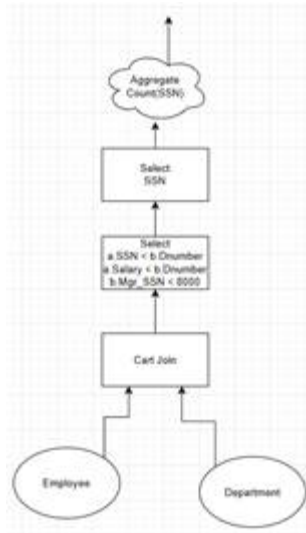
- 1) final DB statistics
- 2) Schema modifications: merged Works_On with Employee, created new Indices (Salary) & (Salary, SSN), decreased size of column data types.
- 3) Memory optimization: increased Buffer size from 65536 B to 65536000 B.
- 4) Index-Tuning: created Indices (Salary) & (Salary,SSN)
- 5) for each query: modifications

Original Query 3:

```
another_complex3 = 'select count(SSN) from Employee a, Department b where a.SSN < b.Dnumber and a.Salary < b.Dnumber and b.Mgr_SSN < 80000'
```

Optimized Query 3:

```
another_complex3_OPT = 'select SSN from (select SSN from Employee) as a inner join (select Dnumber from Department)as b on a.SSN < b.Dnumber inner join (select Dnumber, Mgr_SSN from Department) as d on a.Salary < d.Dnumber where d.Mgr_SSN < 80000'
```

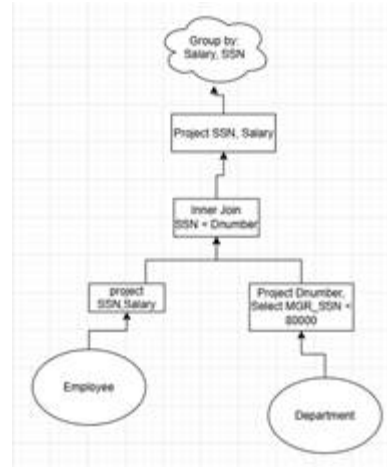





Original Query 4:

```
another_complex4 = 'select SSN , Salary from Employee a, Department b where a.SSN < b.Dnumber and a.Salary < b.Dnumber and b.Mgr_SSN < 80000 group by Salary, SSN;'
```

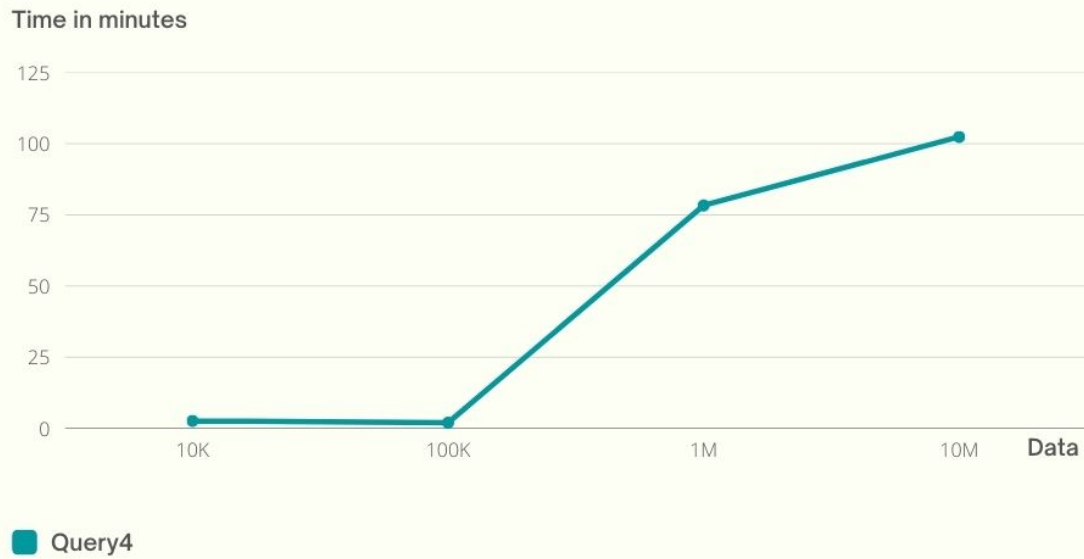
Optimized Query 4:

```
another_complex4_OPT = 'select SSN, Salary from (Select SSN, Salary from Employee)as a inner join (select Dnumber, Mgr_SSN from Department where Mgr_SSN < 80000)as b on a.SSN < b.Dnumber group by Salary, SSN'
```

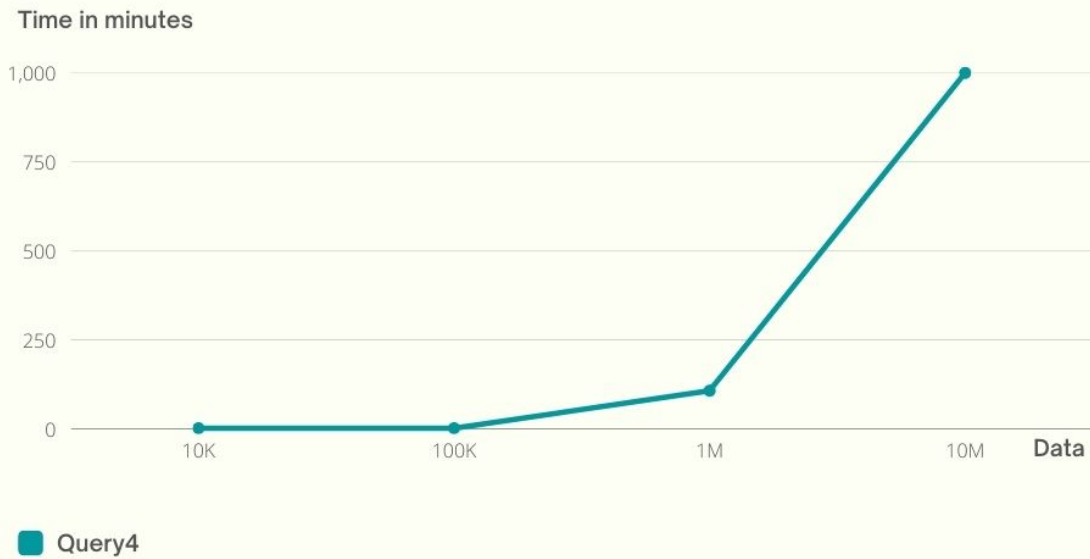




Query 4 without optimizations After Salary Index



Query 4 without optimizations after increasing the buffer size



Query 4

